

# TTIC 31230 Problem Set 3

## Win 2017

Hao Jiang

January 26, 2017

### Problem 1

#### Implementation

Please see attached jupyter notebook for the implementation of Conv / MaxPool / AvePool class as well as the convolutional network architecture.

#### Experiment Result

The experiment result is demonstrated in Figure 1 and I get a test accuracy of 0.3904 after 10 epochs. The execution time on my test machine is around 400s per epoch.

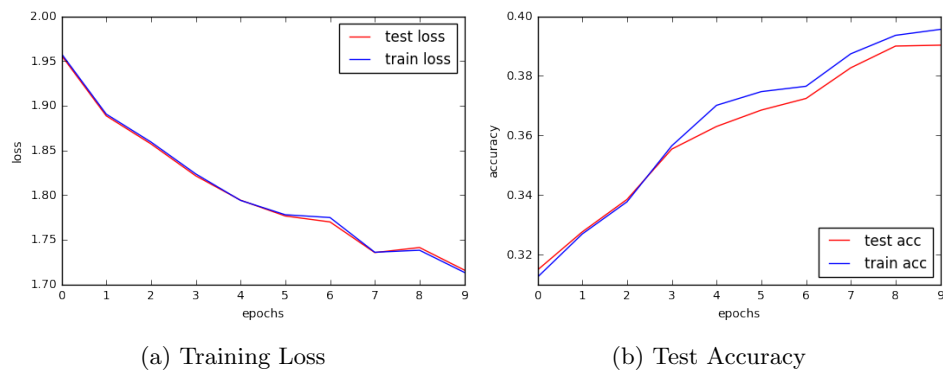


Figure 1: Experiment Result

## Performance Optimization

To speed up the convolution operation, I make heavy use of numpy's ndarray operations instead of writing nested loops. All operations in my implementation only loop on two dimensions, namely the images' width and height.

In the **forward** function of **MaxPool** and **AvePool**, I use numpy's **max** and **average** operation. Let the shape of output value be  $(B, W, H, C)$ , I loop on dimension  $W$  and  $H$ , which indicates a square region on all batch instances and channels. For each such square region, I calculate the max / average on axis 1,2, leaving axis 0 and 3 (which are the index in a batch and channel) unchanged. This operation is efficiently equivalent to compute max / average over each images in a batch and each channels in parallel. It avoids the loop on batch size and number of channels and speed up the operation.

In the **backward** function of **AvePool**, I again only loop on two dimensions. Let the shape of **y.grad** be  $(B, W, H, C)$ . The loop runs on  $W$  and  $H$  and get a ndarray of shape  $(B, C)$  for each  $(W, H)$  pair. For each such ndarray, I use an all-1 square matrix of shape  $(k, k)$  to expand it to shape  $(B, k, k, C)$ , where  $k$  is the size of the square region. Let  $S$  be the stride size, we have

$$\forall b, c, i, j, k_i, k_j, \text{expand}[b, i * S + k_i, j * S + k_j, c] = \text{ygrad}[b, i, j, c]$$

By simply adding all these expanding results together, we will have the gradient to be updated to **x.grad**.

In the **backward** function of **MaxPool**, I use the similar idea. But here in each square region, the gradient will only be backprop to locations holding the maximal value. To implement this, we create a mask by comparing **x.value** to a square matrix with every value equal to the maximal value.

An example is shown in Figure 2. Here the input value is  $\begin{bmatrix} 3 & 6 \\ 7 & 7 \end{bmatrix}$ , and the maximal value is 7. Comparing the input with an all-7 matrix gives us a mask indicating the location of maximal value. When we want backprop a gradient, e.g., 5 to  $X$ , we first expand the gradient to a square and element-wise multiply it with the mask, obtaining the backprop result.

For **Conv**'s **forward** and **backward** method, I follow the hint in problem description and use **numpy.einsum** to manipulate multi-dimensional array multiplications. Again this requires only explicit loop on the  $W$  and  $H$  dimension on the dataset.

$$\begin{aligned}
X &= \begin{bmatrix} 3 & 6 \\ 7 & 7 \end{bmatrix} \\
\left( \begin{bmatrix} 3 & 6 \\ 7 & 7 \end{bmatrix} == \begin{bmatrix} 7 & 7 \\ 7 & 7 \end{bmatrix} \right) &\Rightarrow \begin{bmatrix} 0 & 0 \\ 1 & 1 \end{bmatrix} \\
\text{grad} = 5 &\Rightarrow \\
\begin{bmatrix} 5 & 5 \\ 5 & 5 \end{bmatrix} \circ \begin{bmatrix} 0 & 0 \\ 1 & 1 \end{bmatrix} &\Rightarrow \begin{bmatrix} 0 & 0 \\ 5 & 5 \end{bmatrix}
\end{aligned}$$

Figure 2: Example: Implementing MaxPool.backward